

120??

DOI: *The Canadian Journal of Statistics / La revue canadienne de statistique*

Theorem

[the-

o-

rem]Corollary

[the-

o-

rem]Definition

[theorem]Lemma

[the-

o-

rem]Proposition

[the-

o-

rem]Remark

**Weighted  
Bayesian  
Boot-  
strap  
for  
Scal-  
able  
Bayes**

BlindedA<sup>1\*</sup>  
BlindedB<sup>2</sup>

<sup>1</sup>*Author*

*af-*

*fil-*

*i-*

*a-*

*tions*

*will*

*go*

BLINDED AND BLINDED

*here*

*in*

*the*

*ac-*

*cepted*

*manuscript,*

*but*

*do*

NOT

*in-*

*clude*

*them*

*in*

*your*

*ini-*

*tial*

*sub-*

*mis-*

*sion*

*be-*

*cause*

*it*

*must*

520??

*be*

*anonymous.*

<sup>2</sup>*Second*

*Affiliation*

*Bayesian,*

*Boot-*

*strap,*

*MCMC,*

*Weighted*

*Boot-*

*strap,*

*ABC,*

*Trend*

*Fil-*

*ter-*

*ing,*

*Deep*

*Learn-*

*ing,*

*Ten-*

*sor-*

*Flow,*

*Reg-*

*u-*

BLINDEDA AND BLINDED

lar-

iza-

tion:

MSC

2010:

Pri-

mary

62???

sec-

ondary

62???

We

de-

velop

a

weighted

Bayesian

Boot-

strap

(WBB)

for

ma-

chine

720??

learn-

ing

and

statis-

tics.

WBB

pro-

vides

un-

cer-

tainty

quan-

tifi-

ca-

tion

by

sam-

pling

from

a

high

di-

men-

sional

pos-

te-

rior

dis-

tri-

bu-

tion.

WBB

is

com-

pu-

ta-

tion-

ally

fast

and

scal-

able

us-

ing

only

off-

the-

shelf

op-

ti-



920??

miza-

tion

soft-

ware

such

as

Ten-

sor-

Flow.

We

pro-

vide

reg-

u-

lar-

ity

con-

di-

tions

which

ap-

ply

to

a

wide

range  
of  
ma-  
chine  
learn-  
ing  
and  
sta-  
tis-  
ti-  
cal  
mod-  
els.  
We  
il-  
lus-  
trate  
our  
method-  
ol-  
ogy  
in  
reg-  
u-  
lar-

1120??

ized

re-

gres-

sion,

trend

fil-

ter-

ing

and

deep

learn-

ing.

Fi-

nally,

we

con-

clude

with

di-

rec-

tions

for

fu-

ture

re-

search.

*The*

*Cana-*

*dian*

*Jour-*

*nal*

*of*

*Statis-*

*tics xx: 1–*

25; 20??

© 20?? Sta-

tis-

ti-

cal

So-

ci-

ety

of

Canada

*Résumé:* Insérer

votre

résumé

ici.

We

will

1320??

sup-

ply

a

French

ab-

stract

for

those

au-

thors

who

can't

pre-

pare

it

them-

selves. *La*

*re-*

*vue*

*cana-*

*di-*

*enne*

*de*

*statis-*

*tique* xx: 1–

BLINDED AND BLINDED

25; 20??

© 20?? Société

statis-

tique

du

Canada

---

*\* Author  
to  
whom  
cor-  
re-  
spon-  
dence  
may  
be  
ad-  
dressed.  
E-  
mail:  
In-  
sert  
your  
email  
ad-  
dress  
here  
only  
af-  
ter  
your  
pa-  
per  
has  
been  
ac-  
cepted*

*The Canadian Journal of Statistics / La revue canadienne de statistique* DOI:

1520??

1.

IN-

TRO-

DUC-

TION

Weighted

Bayesian

Boot-

strap

(WBB)

is

a

simulation-

based

al-

go-

rithm

for

as-

sess-

ing

un-

cer-  
tainty  
in  
ma-  
chine  
learn-  
ing  
and  
statis-  
tics.  
Un-  
cer-  
tainty  
quan-  
tifi-  
ca-  
tion  
(UQ)  
is  
an  
ac-  
tive



1720??

area

of

re-

search,

par-

tic-

u-

larly

in

high-

dimensional

in-

fer-

ence

prob-

lems.

Whilst

there

are

com-

pu-

ta-



1920??

con-

texts,

un-

cer-

tainty

as-

sess-

ment

is

still

re-

quired.

De-

vel-

op-

ing

com-

pu-

ta-

tion-

ally

fast

scal-  
able  
al-  
go-  
rithms  
for  
sam-  
pling  
a  
pos-  
te-  
rior  
dis-  
tri-  
bu-  
tion  
is  
a  
no-  
to-  
ri-

2120??

ously

hard

prob-

lem.

WBB

makes

a

con-

tri-

bu-

tion

to

this

lit-

er-

a-

ture

by

show-

ing

how

off-

the-

shelf

op-

ti-

miza-

tion

al-

go-

rithms,

such

as

con-

vex

op-

ti-

miza-

tion

or

stochas-

tic

2320??

gra-

di-

ent

de-

scent

(SGD)

in

Ten-

sor-

Flow

can

also

be

used

to

pro-

vide

un-

cer-

tainty

quan-

tifi-

ca-  
tion.

Our  
work  
builds  
on  
?  
who  
pro-  
vide  
a  
weighted  
like-  
li-  
hood  
Boot-  
strap  
(WLB)  
method  
for  
Bayesian



2520??

in-

fer-

ence.

They

de-

velop

a

weighted

like-

li-

hood

Boot-

strap

al-

go-

rithm

to-

gether

with

the

ap-

pro-

pri-

ate

asymp-

totic

anal-

y-

sis

to

show

that

such

an

al-

go-

rithm

pro-

vides

ef-

fi-

cient

pos-

2720??

te-

rior

sam-

ples.

Their

boot-

strap

pro-

ce-

dure

ex-

ploits

the

fact

that

the

pos-

te-

rior

dis-

tri-

bu-  
tion  
cen-  
tered  
at  
the  
max-  
i-  
mum  
like-  
li-  
hood  
es-  
ti-  
mate  
(MLE)  
has  
a  
sec-  
ond  
or-

2920??

der

ex-

pan-

sion

that

also

de-

pends

on

the

prior

and

its

deriva-

tive.

The

weighted

Bayesian

Boot-

strap

(WBB)

cal-

cu-

lates

a

se-

ries

of

pos-

te-

rior

modes

rather

than

MLEs.

This

has

the

ad-

van-

tage

that

3120??

high

di-

men-

sional

pos-

te-

rior

modes

are

read-

ily

avail-

able

par-

tic-

u-

larly

us-

ing

the

reg-

lar-

ized

es-

ti-

mates

are

fast

to

com-

pute

from

con-

vex

op-

ti-

miza-

tion

meth-

ods

or

stochas-



3320??

tic

gra-

di-

ent

de-

scent

(SGD)

for

neu-

ral

net-

work

ar-

chi-

tec-

tures

such

as

deep

learn-

ing.

By

link-

ing

WLB

and

WBB,

with

modern-

day

op-

ti-

miza-

tion

to

cal-

i-

brate

es-

ti-

mate,

we

3520??

pro-

vide

a

frame-

work

for

un-

cer-

tainty

quan-

tifi-

ca-

tion.

Uncertainty

es-

ti-

mates

are

pro-

vided

at

lit-

tle

to

no

ex-

tra

cost.

Quan-

ti-

fy-

ing

un-

cer-

tainty

is

typ-

i-

cally

un-

avail-

able

3720??

in

a

purely

reg-

u-

lar-

iza-

tion

op-

ti-

miza-

tion

method.

An-

other

fea-

ture

that

is

straight-

for-

ward  
to  
add  
is  
a  
reg-  
u-  
lar-  
iza-  
tion  
path  
across  
hyper-  
parameters.  
This  
is  
so  
much  
eas-  
ier  
than

3920??

tra-

di-

tional

Bayesian

to

do

prior

sen-

si-

tiv-

ity

anal-

y-

sis

where

hyper-

parameters

are

hard

to

as-

sess.

Rather

we

use

pre-

dic-

tive

cross-

validation

tech-

niques.

The

rest

of

the

pa-

per

is

out-

lined

as



4120??

fol-

lows.

Sec-

tion

2

de-

vel-

ops

our

weighted

Bayesian

Boot-

strap

(WBB)

al-

go-

rithm.

Sec-

tion

3

pro-

vides

an

ap-

pli-

ca-

tion

to

high

di-

men-

sional

sparse

re-

gres-

sion,

trend

fil-

ter-

ing

and

deep

4320??

learn-

ing.

WBB

can

also

be

ap-

plied

to

Bayesian

tree

mod-

els

(?).

Fi-

nally,

Sec-

tion

4

con-

cludes

with

di-

rec-

tions

for

fu-

ture

re-

search.

Ar-

eas

for

fu-

ture

study

in-

clude

Boot-

strap

fil-

ters

4520??

in

state-

space

mod-

els

(?)

and

com-

par-

i-

son

with

the

resampling-

sampling

per-

spec-

tive

to

se-

quen-

UNINDEDA AND UNINDEDA

tial

Bayesian

in-

fer-

ence

(?),

etc.

2.

WEIGHTED

BAYESIAN

BOOT-

STRAP

Let

$y$

an

$n$ -

vector

of

out-

comes,

$\theta$

4720??

de-

notes

a

*d*-

dimensional

pa-

ram-

e-

ter

of

in-

ter-

est

and

*A*

a

fixed

$n \times$

*d*

ma-

trix

whose

rows

are

the

de-

sign

points

(or

“fea-

tures”)

$a_i^T$

where

we

in-

dex

ob-

ser-

va-

tions

by

$i$



4920??

and

pa-

ram-

e-

ters

by

$j$ .

A

large

num-

ber

of

ma-

chine

learn-

ing

and

sta-

tis-

ti-

cal

prob-  
lems  
can  
be  
ex-  
pressed  
in  
the  
form

$$\underset{\theta \in \mathcal{R}^d}{\text{minimize}} \quad l(y|\theta) + \lambda \phi(\theta),$$

where

$$l(y|\theta) = \sum_{i=1}^n \log f(y_i; a_i^\top \theta)$$

is

a

mea-

sure

of

fit

(or

5120??

“em-

pir-

i-

cal

risk

func-

tion”)

de-

pend-

ing

im-

plic-

itly

on

*A*

and

*y*.

The

penalty

func-

tion

or

reg-

u-

lar-

iza-

tion

term,

$\lambda\phi(\theta)$ ,

ef-

fects

a

fa-

vor-

able

bias-

variance

trade-

off.

We

al-

low

for

5320??

the

pos-

si-

bil-

ity

that

$\phi(\theta)$

may

have

points

in

its

do-

main

where

it

fails

to

be

dif-

fer-

en-  
tiable.

Suppose

that

we

ob-

serve

data,

$y =$

$(y_1, \dots, y_n)$

from

a

model

pa-

ram-

e-

ter-

ized

by

$\theta$ .

For

5520??

ex-

am-

ple,

we

might

have

a

prob-

a-

bilis-

tic

model

that

de-

pends

on

a

pa-

ram-

e-

ter,

$\theta$ ,

where

$$p(y|\theta)$$

is

known

as

the

like-

li-

hood

func-

tion.

Equiv-

a-

lently,

we

can

de-

fine

a

mea-

sure



5720??

of

fit

$$l(y|\theta) =$$

$$\log f(y; \theta) =$$

$$\log p(y|\theta).$$

We

will

make

use

of

the

fol-

low-

ing

(i)

Let

$$\hat{\theta}_n :=$$

$$\operatorname{argmax}_{\theta} p(y|\theta)$$

be

the

MLE,

(ii)

Let

$$\theta_n^* :=$$

$$\operatorname{argmax}_{\theta} p(\theta|y)$$

be

the

pos-

te-

rior

mode,

(iii)

Let

$$\bar{\theta}_n :=$$

$$E(\theta|y)$$

be

the

pos-

te-

rior

mean.

5920??

We

now

de-

velop

a

key

du-

al-

ity

be-

tween

reg-

u-

lar-

iza-

tion

and

pos-

te-

rior

boot-

strap

sim-

u-

la-

tion.

2.1.

Bayesian

Reg-

u-

lar-

iza-

tion

Du-

al-

ity

From

the

Bayesian

per-

spec-

tive,

the

6120??

mea-

sure

of

fit,

$l(y|\theta) =$

$-\log f(y; \theta),$

and

the

penalty

func-

tion,

$\lambda\phi(\theta),$

cor-

re-

spond

to

the

neg-

a-

tive

log-

a-

rithms

of

the

like-

li-

hood

and

prior

dis-

tri-

bu-

tion

in

the

hi-

er-

ar-

chi-

cal

6320??

model

$$f(y; \theta) = p(y|\theta) \propto \exp\{-l(y|\theta)\}, \quad p(\theta) \propto \exp\{-\lambda\phi(\theta)\}$$

$$p(\theta|y) \propto \exp\{-(l(y|\theta) + \lambda\phi(\theta))\}.$$

The

prior

is

not

nec-

es-

sar-

ily

proper

but

the

pos-

te-

rior,

$$p(\theta|y) \propto$$

$$p(y|\theta)p(\theta),$$

may

still

be

proper.

This

pro-

vides

an

equiv-

a-

lence

be-

tween

reg-

u-

lar-

iza-

tion

and

Bayesian

meth-

ods.



6520??

For

ex-

am-

ple,

re-

gres-

sion

with

a

least

squares

log-

likelihood

sub-

ject

to

a

penalty

such

as

an

$L^2$ -

norm

(ridge)

Gaus-

sian

prob-

a-

bil-

ity

model

or

$L^1$ -

norm

(lasso)

dou-

ble

ex-

po-

nen-

tial

prob-

6720??

a-

bil-

ity

model.

We

then

have

$$\hat{\theta}_n = \underset{\theta \in \Theta}{\operatorname{argmax}} l(y|\theta), \quad (2)$$

$$\theta_n^* = \underset{\theta \in \Theta}{\operatorname{argmax}} \{l(y|\theta) + \lambda \phi(\theta)\} \quad (3)$$

Let

$\partial$

be

the

sub-

d-

if-

fer-

en-

tial

op-

er-

a-

tor.

Then

a

nec-

es-

sary

and

suf-

fi-

cient

con-

di-

tion

for

$\theta^*$

to

min-

i-

mize

$l(y|\theta) +$

6920??

$\lambda\phi(\theta)$

is

$$0 \in \partial \{l(y|\theta) + \lambda\phi(\theta)\} = \nabla l(y|\theta) + \lambda\partial\phi(\theta)$$

the

sum

of

a

point

and

a

set.

The

op-

ti-

miza-

tion

lit-

er-

a-

ture

char-

ac-

ter-

izes

 $\theta^*$ 

as

the

fixed

point

of

a

prox-

i-

mal

op-

er-

a-

tor

$$\theta^* =$$
$$\text{prox}_{\gamma\phi}\{\theta^* -$$
$$\lambda \nabla f(\theta^*)\},$$

7120??

see

?

and

?

for

fur-

ther

dis-

cus-

sion.

A

gen-

eral

class

of

nat-

u-

ral

ex-

po-

nen-

tial  
fam-  
ily  
mod-  
els  
can  
be  
ex-  
pressed  
in  
terms  
of  
the  
Breg-  
man  
di-  
ver-  
gence  
of  
the  
dual



7320??

of

the

cu-

mu-

lant

trans-

form.

Let

$\phi$

be

the

con-

ju-

gate

Leg-

en-

dre

trans-

form

of

$\psi$ .

Hence

$$\psi(\theta) = \sup_{\mu} (\mu^{\top} \theta - \phi(\mu)).$$

Then

we

can

write

$$\begin{aligned} p_{\psi}(y|\theta) &= \exp(y^{\top} \theta - \psi(\theta) - h_{\psi}(y)) \\ &= \exp \left\{ \inf_{\mu} ((y - \mu)^{\top} \theta - \phi(\mu)) - h_{\psi}(y) \right\} \\ &= \exp(-D_{\phi}(y, \mu(\theta)) - h_{\phi}(y)) \end{aligned}$$

where

the

in-

fi-

mum

is

at-

tained

at

7520??

$$\mu(\theta) =$$

$$\phi'(\theta)$$

is

the

mean

of

the

ex-

po-

nen-

tial

fam-

ily

dis-

tri-

bu-

tion.

We

rewrite

$$h_{\psi}(y)$$

is

terms

of

the

cor-

rec-

tion

term

and

$h_\phi(y)$ .

Here

there

is

a

du-

al-

ity

as

$D_\phi$

can

be

in-

7720??

ter-

preted

as

a

Breg-

man

di-

ver-

gence.

For

a

wide

range

of

non-

smooth

ob-

jec-

tive

func-

tions/statistical

mod-

els,

re-

cent

reg-

u-

lar-

iza-

tion

meth-

ods

pro-

vide

fast,

scal-

able

al-

go-

rithms

for

7920??

cal-

cu-

lat-

ing

es-

ti-

mates

of

the

form

(??),

which

can

also

be

viewed

as

the

pos-

te-

rior

mode.

There-

fore

as

 $\lambda$ 

varies

we

ob-

tain

a

full

reg-

u-

lar-

iza-

tion

path

as

a

form

of



8120??

prior

sen-

si-

tiv-

ity

anal-

y-

sis.

?

and

?

con-

sid-

ered

sce-

nar-

ios

where

pos-

te-

rior

modes

can

be

used

as

pos-

te-

rior

means

from

aug-

mented

prob-

a-

bil-

ity

mod-

els.

More-

over,

8320??

in

their

orig-

i-

nal

foun-

da-

tion

of

the

Weighted

Like-

li-

hood

Boot-

strap

(WLB),

?

in-

tro-

duced

the  
con-  
cept  
of  
the  
im-  
plicit  
prior.  
Clearly  
this  
is  
an  
av-  
enue  
for  
fu-  
ture  
re-  
search.

8520??

2.2.

WBB

Al-

go-

rithm

We

now

de-

fine

the

weighted

Bayesian

Boot-

strap

(WBB).

Fol-

low-

ing

?,

we

con-

struct

a

ran-

domly

weighted

pos-

te-

rior

dis-

tri-

bu-

tion

de-

noted

by

$$\mathbf{w} = (w_1, \dots, w_n, w_p), p_{\mathbf{w}}(\theta|y) \propto \prod_{i=1}^n p(y_i|\theta)^{w_i} p(\theta)^{w_p}$$

where

the

weights

$$w_p, w_i \sim$$

$$Exp(1)$$

8720??

are

ran-

domly

gen-

er-

ated

weights.

It's

equiv-

a-

lent

to

draw

$w_i =$

$\log(1/U_i)$

where

$U_i$ 's

are

i.i.d.

Uni-

form

$(0,1)$ ,

which

is

mo-

ti-

vated

by

the

uni-

form

Dirich-

let

dis-

tri-

bu-

tion

for

multi-

no-

mial

data.

We



8920??

have

used

the

fact

that

for

i.i.d.

ob-

ser-

va-

tions,

the

like-

li-

hood

can

be

fac-

tor-

ized

as

$$p(y|\theta) =$$

$$\prod_{i=1}^n p(y_i|\theta).$$

This

is

not

cru-

cial

for

our

anal-

y-

sis

but

is

a

com-

mon

as-

sump-

tion.

Let

9120??

$\theta_{\mathbf{w},n}^*$

de-

note

the

mode

of

this

reg-

u-

lar-

ized

dis-

tri-

bu-

tion.

Again,

there

is

an

equiv-

a-

lence

$$\theta_{\mathbf{w},n}^* := \arg \min_{\theta} p_{\mathbf{w}}(\theta|y) \equiv \arg \min_{\theta} \sum_{i=1}^n w_i l_i(y_i|\theta) + \lambda w_p \phi(\theta)$$

where

$$l_i(y_i|\theta) =$$

$$-\log p(y_i|\theta)$$

and

$$\lambda \phi(\theta) =$$

$$-\log p(\theta).$$

Note

that

we

have

a

weighted

like-

li-

hood

and

a

9320??

new

reg-

u-

lar-

iza-

tion

pa-

ram-

e-

ter,

$\lambda w_p$ .

The

crux

of

our

pro-

ce-

dure

is

to

cre-

ate

a

sam-

ple

of

the

weighted

pos-

te-

rior

modes

$\{\theta_{w,n}^*\}$

(com-

pu-

ta-

tion-

ally

cheap

as

each

sub-

9520??

problem

can

be

solved

via

op-

ti-

miza-

tion).

Our

main

re-

sult

is

the

fol-

low-

ing:

**Algorithm:**

**Weighted**

## Bayesian

### Boot-

### strap

### (WBB)

1.

Iterate:

sam-

ple

$\mathbf{w} =$

$\{w_1, w_2, \dots, w_n, w_p\}$

via

ex-

po-

nen-

tials.

$w_p, w_i \sim$

$Exp(1)$ .

2.

For

each

$\mathbf{w},$



9720??

solve

$$\theta_{\mathbf{w},n}^* =$$

$$\sum_{i=1}^n w_i l_i(\theta) +$$

$$\lambda w_p \phi(\theta).$$

The

WBB

al-

go-

rithm

is

fast

and

scal-

able

to

com-

pute

a

reg-

u-

lar-

ized

es-

ti-

ma-

tor.

For

a

large

num-

ber

of

pop-

u-

lar

pri-

ors,

the

min-

i-

miz-

ing

9920??

so-

lu-

tion

$\theta_{\mathbf{w},n}^*$

in

the

sec-

ond

step

can

be

di-

rectly

ob-

tained

via

reg-

u-

lar-

iza-

tion

pack-

ages

such

as

glmnet

by

Trevor

Hastie

and

genlasso

by

Tay-

lor

Arnold.

When

the

like-

li-

hood

func-

tion

or

the

10120??

prior  
is  
spe-  
cially  
de-  
signed,  
Stochas-  
tic  
Gra-  
di-  
ent  
De-  
scent  
(SGD)  
is  
pow-  
er-  
ful  
and  
fast  
enough

to  
solve  
the  
min-  
i-  
miza-  
tion  
prob-  
lem.  
It  
can  
be  
eas-  
ily  
im-  
ple-  
mented  
in  
Ten-  
sor-  
Flow

10320??

once

the

ob-

jec-

tive

func-

tion

is

spec-

i-

fied.

See

Ap-

pendix

(??)

and

?

for

fur-

ther

dis-

cus-

sion.

The

next

sec-

tion

builds

on

?

and

de-

rives

asympt-

otic

prop-

er-

ties

of

the

weighted

Bayesian



10520??

Boot-

strap.

We

sim-

ply

add

the

reg-

u-

lar-

ized

fac-

tor.

To

choose

the

amount

of

reg-

u-

lar-

iza-

tion

$\lambda$ ,

we

can

use

the

marginal

like-

li-

hood

$m_\lambda(y)$ ,

es-

ti-

mated

by

bridge

sam-

pling

(?)

or

10720??

sim-

ply

us-

ing

pre-

dic-

tive

cross-

validation.

2.3.

WBB

Prop-

er-

ties

The

fol-

low-

ing

propo-

si-

tion

which

fol-

lows

from

the

The-

o-

rem

2

in

?

sum-

maries

the

prop-

er-

ties

of

WBB.

**Proposition***The**weighted**Bayesian**Boot-**strap**draws**are**ap-**prox-**i-**mate**pos-**te-**rior**sam-**ples*

$$\{\theta_{\mathbf{w},n}^{*(k)}\}_{k=1}^K \sim p(\theta|y).$$



11120??

$$e^{-nl_n(\theta)}p(\theta)$$

for

suf-

fi-

ciently

large

$n$

will

be

dom-

i-

nated

by

the

like-

li-

hood

term.

Ex-

pand-

ing

$l_n(\theta)$

around

its

max-

i-

mum,

$\hat{\theta}$ ,

and

defin-

ing

$J_n(\hat{\theta}) =$

$nj(\hat{\theta})$

as

the

ob-

served

in-

for-

ma-

tion

ma-



11320??

trix

gives

the

tra-

di-

tional

nor-

mal

ap-

prox-

i-

ma-

tion

for

the

pos-

te-

rior

dis-

tri-

bu-

tion

$$\theta \sim N_d \left( \hat{\theta}_n, J_n^{-1}(\hat{\theta}) \right)$$

where

$$\hat{\theta}_n$$

is

the

MLE.

A

more

ac-

cu-

rate

ap-

prox-

i-

ma-

tion

is

ob-

tained

11520??

by

ex-

pand-

ing

around

the

pos-

te-

rior

mode,

$\theta^*$ ,

which

we

will

ex-

exploit

in

our

weighted

Bayesian

Boot-

strap.

Now

we

have

the

asymptotic

dis-

tri-

bu-

tional

ap-

prox-

i-

ma-

tion

$$\theta \sim N_d(\theta^*, J_n^{-1}(\theta^*))$$

where

$$\theta_n^* :=$$

$$\arg \max_{\theta} p(\theta|y)$$

11720??

is

the

pos-

te-

rior

mode.

The

use

of

the

pos-

te-

rior

mode

here

is

cru-

cially

im-

por-

tant  
as  
it's  
the  
mode  
that  
is  
com-  
pu-  
ta-  
tion-  
ally  
avail-  
able  
from  
Ten-  
sor-  
Flow  
and  
Keras.  
Ap-

11920??

prox-

i-

mate

nor-

mal-

ity

and

sec-

ond

or-

der

ap-

prox-

i-

ma-

tion

also

holds,

see

?,

?

and

?

for

fu-

ture

dis-

cus-

sion.

Specif-

i-

cally,

$$\sqrt{nI(\hat{\theta}_n)} \left( \theta_n^* - \hat{\theta}_n \right) \stackrel{D}{=} Z$$

where

$$Z \sim$$

$$N(0, 1)$$

is

a

stan-

dard

Nor-

mal



12120??

vari-

able.

The

con-

di-

tional

pos-

te-

rior

sat-

is-

fies

$\mathbb{P}\left(|\theta_n^* - \hat{\theta}_n| > \epsilon\right) \rightarrow$

0

for

each

$\epsilon >$

0

as

$n \rightarrow$

$\infty$ .

In  
the  
'large  
 $p'$   
case,  
a  
num-  
ber  
of  
re-  
sults  
are  
avail-  
able  
for  
pos-  
te-  
rior  
con-  
cen-  
tra-

12320??

tion,

for

ex-

am-

ple,

see

?

for

sparse

high

di-

men-

sional

mod-

els.

3.

AP-

PLI-

CA-

TIONS

Consider

now

a

num-

ber

of

sce-

nar-

ios

to

as-

sess

when

WBB

cor-

re-

sponds

12520??

to

a

full

Bayesian

pos-

te-

rior

dis-

tri-

bu-

tion.

3.1.

Lasso

First,

a

sim-

ple

uni-

vari-

ate

nor-

mal  
means  
prob-  
lem  
with  
a  
lasso  
prior  
where

$$y|\theta \sim N(\theta, 1^2), \quad \theta \sim Laplace(0, 1/\lambda)$$

Given  
the  
i.i.d.  
ex-  
po-  
nen-  
tial  
weights  
 $w_1$   
and

12720??

$w_2$ ,

the

weighted

pos-

te-

rior

mode

$\theta_{\mathbf{w}}^*$

is

$$\theta_{\mathbf{w}}^* = \underset{\theta \in \Theta}{\operatorname{arg\,min}} \left\{ \frac{w_1}{2} (y - \theta)^2 + \lambda w_2 |\theta| \right\}.$$

This

is

suf-

fi-

ciently

sim-

ple

for

an

ex-

act

WBB

so-

lu-

tion

in

terms

of

soft

thresh-

old-

ing:

$$\theta_{\mathbf{w}}^* = \begin{cases} y - \lambda w_2/w_1 & \text{if } y > \lambda w_2/w_1, \\ y + \lambda w_2/w_1 & \text{if } y < -\lambda w_2/w_1, \\ 0 & \text{if } |y| \leq \lambda w_2/w_1. \end{cases}$$

The

WBB

mean

$$E_{\mathbf{w}}(\theta_{\mathbf{w}}^*|y)$$

is



12920??

ap-

prox-

i-

mated

by

the

sam-

ple

mean

of

$\{\theta_{\mathbf{w}}^{*(k)}\}_{k=1}^K$ .

On

the

other

hand,

?

gives

the

ex-

pres-

sion

for

the  
pos-  
te-  
rior  
mean,

$$\begin{aligned} E(\theta|y) &= \frac{\int_{-\infty}^{\infty} \theta \exp \{-(y-\theta)^2/2 - \lambda|\theta|\} d\theta}{\int_{-\infty}^{\infty} \exp \{-(y-\theta)^2/2 - \lambda|\theta|\} d\theta} \\ &= \frac{F(y)}{F(y) + F(-y)}(y + \lambda) + \frac{F(-y)}{F(y) + F(-y)}(y - \lambda) \\ &= y + \frac{F(y) - F(-y)}{F(y) + F(-y)}\lambda \end{aligned}$$

where

$$F(y) =$$

$$\exp(y)\Phi(-y -$$

$$\lambda)$$

and

$$\Phi(\cdot)$$

is

the

c.d.f.

of

13120??

stan-

dard

nor-

mal

dis-

tri-

bu-

tion.

We

plot

the

WBB

mean

ver-

sus

the

ex-

act

pos-

te-

rior

mean

in

Fig-

ure

(??).

In-

ter-

est-

ingly,

WBB

al-

go-

rithm

gives

3.2.

sparser

Di-

pos-

a-

tes

Data

Means.

il-  
simple.pdf

lus

trate

13320??

our

method-

ol-

ogy,

we

use

weighted

Bayesian

Boot-

strap

(WBB)

on

the

clas-

sic

di-

a-

betes

dataset.

The

mea-

sure-

ments

for

442

di-

a-

betes

pa-

tients

are

ob-

tained

( $n =$

442),

with

10

base-

line

vari-

ables

( $p =$

10),

13520??

such

as

age,

sex,

body

mass

in-

dex,

av-

er-

age

blood

pres-

sure,

and

six

blood

serum

mea-

sure-

ments.

The  
like-  
li-  
hood  
func-  
tion  
is  
given  
by

$$l(y|\beta) = \prod_{i=1}^n p(y_i|\beta)$$

where

$$p(y_i|\beta) = \frac{1}{\sqrt{2\pi}\sigma} \exp \left\{ -\frac{1}{2\sigma^2} (y_i - x_i'\beta)^2 \right\}.$$

We  
draw  
1000  
sets  
of  
weights



13720??

$\mathbf{w} =$

$\{w_i\}_{i=1}^{n+1}$

where

$w_i$ 's

are

i.i.d.

ex-

po-

nen-

tials.

For

each

weight

set,

the

weighted

Bayesian

es-

ti-

mate

$\beta_{\mathbf{w}}^*$

is

cal-  
cu-  
lated  
us-  
ing  
(??)  
via  
the  
reg-  
u-  
lar-  
iza-  
tion  
method  
in  
the  
pack-  
age  
glmnet.

$$\hat{\beta}_{\mathbf{w}} := \arg \min_{\beta} \sum_{i=1}^n w_i (y_i - x_i' \beta)^2 + \lambda w_{n+1} \sum_{j=1}^p |\beta_j| .$$

The

reg-

u-

lar-

iza-

tion

fac-

tor

$\lambda$

is

cho-

sen

by

cross-

validation

with

un-

weighted

like-

li-

hood.

The

weighted

Bayesian

Boot-

strap

is

also

per-

formed

with

fixed

prior,

namely,

$w_{n+1}$

is

set

to

be

14120??

1

for

all

boot-

strap

sam-

ples.

?

an-

a-

lyze

the

same

dataset

us-

ing

the

Bayesian

Bridge

es-

ti-

ma-  
tor  
and  
sug-  
gest  
MCMC  
sam-  
pling  
from  
the  
pos-  
te-  
rior.

To  
com-  
pare  
our  
WBB  
re-  
sults  
we

14320??

also

run

the

Bayesian

bridge

es-

ti-

ma-

tion.

Here

the

Bayesian

set-

ting

we

use

is

$p(\beta, \sigma^2) = p(\beta|\sigma^2)p(\sigma^2)$ , where  $p(\sigma^2) \propto 1/\sigma^2$ .

The

prior

on

$\beta$ ,

with

suit-

able

nor-

mal-

iza-

tion

con-

stant

$C_\alpha$ ,

is

given

by

$$p(\beta) = C_\alpha \exp\left(-\sum_{j=1}^p |\beta_j/\tau|^\alpha\right).$$

The

hyper-

parameter

is



drawn

as

$\nu =$

$\tau^{-\alpha} \sim$

Figure 2,

(21) where

shows

the

re-

sults

of diabetes.pdf  
all

these

three

meth-

ods

(the

weighted

Bayesian

Boot-

strap

with

FIGURE 2:  
Diabetes  
ex-  
am-  
ple:  
the  
the  
weighted  
Bayesian  
Boot-  
strap  
(with  
fixed  
prior  
and  
weighted  
prior)  
and  
Bayesian  
Bridge  
are  
used  
to  
draw  
from  
the  
marginal  
pos-  
te-  
ri-  
ors  
for  
 $\beta_j$ 's,  
j  
=  
1,2,...10.

fixed  
prior  
/  
weighted  
prior  
and  
the  
Bayesian  
Bridge).  
Marginal  
pos-  
te-  
ri-  
ors  
for  
 $\beta_j$ 's  
are  
pre-  
sented.  
One  
no-

14720??

table

fea-

ture

is

that

the

weighted

Bayesian

Boot-

strap

tends

to

in-

tro-

duce

more

spar-

sity

than

Bayesian

Bridge

does.

For

ex-

am-

ple,

the

weighted

Bayesian

Boot-

strap

pos-

te-

ri-

ors

of

age,

ldl

and

tch

have

higher

14920??

spikes

lo-

cated

around

0,

com-

pared

with

the

Bayesian

Bridge

ones.

For

$t_c$ ,

$h_{dl}$ ,

$t_{ch}$

and

$glu$ ,

multi-

modes

in

the  
marginal  
pos-  
te-  
ri-  
ors  
are  
ob-  
served.

In  
gen-  
eral,  
the  
pos-  
te-  
ri-  
ors  
with  
fixed  
pri-  
ors

15120??

are

more

con-

cen-

trated

than

those

with

ran-

domly

weighted

pri-

ors.

This

dif-

fer-

ence

is

nat-

u-

rally

at-

tributed

to

the

cer-

tainty

in

the

prior

weights.

3.3.

Trend

Fil-

ter-

ing

The

gen-

er-

al-

ized

lasso

solves

the



15320??

op-

ti-

miza-

tion

prob-

lem:

$$\beta^* = \arg \min_{\beta} \{l(y|\beta) + \lambda \phi(\beta)\} \quad (6)$$

$$= \arg \min_{\beta} \frac{1}{2} \|y - X\beta\|_2^2 + \lambda \|D\beta\|_1 \quad (7)$$

where

$$l(y|\beta) =$$

$$\frac{1}{2} \|y -$$

$$X\beta\|_2^2$$

is

the

neg-

a-

tive

log-

likelihood.

$$D \in$$

$\mathcal{R}^{m \times p}$

is

a

penalty

ma-

trix

and

$\lambda\phi(\beta) =$

$\lambda\|D\beta\|_1$

is

the

neg-

a-

tive

log-

prior

or

reg-

u-

lar-

iza-

15520??

tion

penalty.

There

are

fast

path

al-

go-

rithms

for

solv-

ing

this

prob-

lem

(see

genlasso

pack-

age).

As

a

sub-

prob-

lem,

poly-

no-

mial

trend

fil-

ter-

ing

(??)

is

re-

cently

in-

tro-

duced

for

piece-

15720??

wise

poly-

no-

mial

curve-

fitting,

where

the

knots

and

the

pa-

ram-

e-

ters

are

cho-

sen

adap-

tively.

In-

tu-  
itively,  
the  
trend-  
filtering  
es-  
ti-  
ma-  
tor  
is  
sim-  
i-  
lar  
to  
an  
adap-  
tive  
spline  
model:  
it  
pe-

15920??

nal-

izes

the

dis-

crete

deriva-

tive

of

or-

der

$k$ ,

re-

sult-

ing

in

piece-

wise

poly-

no-

mi-

als

of  
higher  
de-  
gree  
for  
larger  
 $k$ .

Specifically,

$$X =$$

$$I_p$$

in  
the  
trend  
fil-  
ter-  
ing  
set-  
ting  
and  
the



16120??

data

$y =$

$(y_1, \dots, y_p)$

are

as-

sumed

to

be

mean-

ing-

fully

or-

dered

from

1

to

$p$ .

The

penalty

ma-

trix

is  
 spe-  
 cially  
 de-  
 signed  
 by  
 the  
 dis-  
 crete  
 $(k +$   
 $1)-$   
 th  
 or-  
 der  
 deriva-  
 tive,

$$D^{(1)} = \begin{bmatrix} -1 & 1 & 0 & \dots & 0 & 0 \\ 0 & -1 & 1 & \dots & 0 & 0 \\ \dots & \dots & \dots & \dots & \dots & \dots \\ 0 & 0 & 0 & \dots & -1 & 1 \end{bmatrix}_{(p-1) \times p}$$

16320??

and

$$D^{(k+1)} =$$

$$D^{(1)} D^{(k)}$$

for

$$k =$$

$$1, 2, 3, \dots$$

For

ex-

am-

ple,

the

log-

prior

in

lin-

ear

trend

fil-

ter-

ing

is

ex-

plic-

itly

writ-

ten

as

$$\lambda \sum_{i=1}^{p-2} |\beta_{i+2} -$$

$$2\beta_{i+1} +$$

$$\beta_i|.$$

For

a

gen-

eral

or-

der

$$k >$$

$$1,$$

$$\|D^{(k+1)}\beta\|_1 = \sum_{i=1}^{p-k-1} \left| \sum_{j=i}^{i+k+1} (-1)^{(j-i)} \binom{k+1}{j-i} \beta_j \right|.$$

WBB

solves

16520??

the  
fol-  
low-  
ing  
gen-  
er-  
al-  
ized  
lasso  
prob-  
lem  
in  
each  
draw:

$$\begin{aligned}\beta_{\mathbf{w}}^* &=_{\beta} \frac{1}{2} \sum_{i=1}^p w_i (y_i - \beta_i)^2 + \lambda w_{p+1} \|D^{(k)} \beta\|_1 \\ &=_{\beta} \frac{1}{2} \|Wy - W\beta\|_2^2 + \lambda \|D^{(k)} \beta\|_1 \\ &= W^{-1} \frac{1}{\tilde{\beta}} \|\tilde{y}_{\mathbf{w}} - \tilde{\beta}_{\mathbf{w}}\|_2^2 + \lambda \|\tilde{D}_{\mathbf{w}}^{(k)} \tilde{\beta}_{\mathbf{w}}\|_1\end{aligned}$$

where

$$W = \text{diag}(\sqrt{w_i}/\sqrt{w_{p+1}}, \dots, \sqrt{w_p}/\sqrt{w_{p+1}})$$

and

$$\tilde{y}_{\mathbf{w}} = Wy, \tilde{\beta}_{\mathbf{w}} = W\beta, \tilde{D}_{\mathbf{w}}^{(k)} = D^{(k)}W^{-1}.$$

To

il-

lus-

trate

our

method,

we

sim-

u-

late

data

$y_i$

from

a

16720??

Fourier

se-

ries

re-

gres-

sion

$$y_i = \sin\left(\frac{4\pi}{500}i\right) \exp\left(\frac{3}{500}i\right) + \epsilon_i$$

for

$i =$

1, 2, ..., 500,

where

$\epsilon_i \sim$

$N(0, 2^2)$

are

i.i.d.

Gaus-

sian

noises.

The

cu-

bic  
trend  
fil-  
ter-  
ing  
re-  
sult  
is  
given  
in  
Fig-  
ure  
(??).

For  
each  
 $i$ ,  
the  
weighted  
Bayesian  
Boot-



16920??

strap

gives

a

group

of

es-

ti-

mates

$\{\beta_{\mathbf{w}}^*(i)\}_{j=1}^T$

where

$T$

is

the

to-

tal

num-

ber

of

draws.

The

stan-

dard

er-

ror

of

$\hat{\beta}_i$

is

eas-

ily

com-

puted

us-

ing

these

weighted

boot-

strap

es-

ti-

mates.

17120??

3.4.

Deep

Learn-

ing:

MNIST

Ex-

am-

ple

Deep

learn-

ing

is

a

form

of

ma-

chine

learn-

ing

that

uses

hi-

er-

ar-

chi-

cal

ab-

tract

lay-

ers

of

1a-

tent

vari-

ables

to

per-

form

pat-

tern

match-

ing

and

pre-

17320??

dic-

tion.

?

take

a

Bayesian

prob-

a-

bilis-

tic

per-

spec-

tive

and

pro-

vide

a

num-

ber

of

in-

sights

into

more

ef-

fi-

cient

al-

go-

rithms

for

op-

ti-

miza-

tion

and

hyper-

parameter

tun-

ing.

17520??

The  
gen-  
eral  
goal  
is  
to  
finds  
a  
pre-  
dic-  
tor  
of  
an  
out-  
put  
 $y$   
given  
a  
high  
di-  
men-

sional

in-

put

$x$ .

For

a

clas-

si-

fi-

ca-

tion

prob-

lem,

$y \in$

$\{1, 2, \dots, K\}$

is

a

dis-

crete

vari-

able



17720??

and

can

be

coded

as

a

$K$ -

dimensional

0-

1

vec-

tor.

The

model

is

as

fol-

lows.

Let

$z^{(l)}$

de-

note

the

$l$ -

th

layer,

and

so

$x =$

$z^{(0)}$ .

The

fi-

nal

out-

put

is

the

re-

sponse

$y$ ,

which

can

17920??

be

nu-

meric

or

cat-

e-

gor-

i-

cal.

A

deep

pre-

dic-

tion

rule

is

then

$$\begin{aligned} z^{(1)} &= f^{(1)}\left(W^{(0)}x + b^{(0)}\right), \\ z^{(2)} &= f^{(2)}\left(W^{(1)}z^{(1)} + b^{(1)}\right), \\ &\dots \\ z^{(L)} &= f^{(L)}\left(W^{(L-1)}z^{(L-1)} + b^{(L-1)}\right), \\ \hat{y}(x) &= z^{(L)}. \end{aligned}$$

Here,

$W^{(l)}$

are

weight

ma-

tri-

ces,

and

$b^{(l)}$

are

thresh-

old

or

18120??

ac-

ti-

va-

tion

lev-

els.

$f^{(l)}$

is

the

ac-

ti-

va-

tion

func-

tion.

Prob-

a-

bilis-

ti-

cally,

the

out-

put

$y$

in

a

clas-

si-

fi-

ca-

tion

prob-

lem

is

gen-

er-

ated

by

a

prob-

a-

bil-

ity

model

$$p(y|x, W, b) \propto \exp\{-l(y|x, W, b)\}$$

where

$$l(y|x, W, b) =$$

$$\sum_{i=1}^n l_i(y_i|x_i, W, b)$$

is

the

neg-

a-

tive

cross-

entropy,

$$l_i(y_i|x_i, W, b) = l_i(y_i, \hat{y}(x_i)) = \sum_{k=1}^K y_{ik} \log \hat{y}_k(x_i)$$

where

$$y_{ik}$$

is

0

or

1

and

$K =$

10.

Adding

the

neg-

a-

tive

log-

prior

$\lambda\phi(W, b),$

the

ob-

jec-

tive

func-

tion

(neg-

a-

tive

log-



18520??

posterior)

to

be

min-

i-

mized

by

stochas-

tic

gra-

di-

ent

de-

scent

is

$$\mathcal{L}_\lambda(y, \hat{y}) = \sum_{i=1}^n l_i(y_i, \hat{y}(x_i)) + \lambda \phi(W, b).$$

Accordingly,

with

each

draw

of  
weights  
 $w$ ,  
WBB  
pro-  
vides  
the  
es-  
ti-  
mates  
 $(W_w^*, b_w^*)$   
by  
solv-  
ing  
the  
fol-  
low-  
ing  
op-  
ti-  
miza-

18720??

tion

prob-

lem.

$$(W_{\mathbf{w}}^*, b_{\mathbf{w}}^*) =_{W,b} \sum_{i=1}^n w_i l_i(y_i | x_i, W, b) + \lambda w_p \phi(W, b)$$

We

take

the

clas-

sic

MNIST

ex-

am-

ple

to

il-

lus-

trate

the

ap-

pli-

ca-  
tion  
of  
WBB  
in  
deep  
learn-  
ing.  
The  
MNIST  
database  
of  
hand-  
writ-  
ten  
dig-  
its,  
avail-  
able  
from  
Yann  
Le-

18920??

Cun's

web-

site,

has

60,000

train-

ing

ex-

am-

ples

and

10,000

test

ex-

am-

ples.

Here

the

high-

dimensional

$x$

is

a

nor-

mal-

ized

and

cen-

tered

fixed-

size

$(28 \times$

28)

im-

age

and

the

out-

put

$\hat{y}$

is

a

19120??

10-

dimensional

vec-

tor,

where

$i$ -

th

co-

or-

di-

nate

cor-

re-

sponds

to

the

prob-

a-

bil-

ity

of

that

im-

age

be-

ing

the

*i*-

th

digit.

For

sim-

plic-

ity,

we

build

a

2-

layer

neu-

ral



19320??

net-

work

with

layer

sizes

128

and

64

re-

spec-

tively.

There-

fore,

the

di-

men-

sions

of

pa-

ram-

e-

ters  
 are

$$W^{(0)} \in \mathcal{R}^{128 \times 784}, b^{(0)} \in \mathcal{R}^{128},$$

$$W^{(1)} \in \mathcal{R}^{64 \times 128}, b^{(1)} \in \mathcal{R}^{64},$$

$$W^{(2)} \in \mathcal{R}^{10 \times 64}, b^{(0)} \in \mathcal{R}^{10}.$$

The  
 ac-  
 ti-  
 va-  
 tion  
 func-  
 tion  
 $f^{(i)}$   
 is  
 ReLU,  
 $f(x) =$   
 $\max\{0, x\},$   
 and  
 the  
 neg-

19520??

a-

tive

log-

prior

is

spec-

i-

fied

as

$$\lambda\phi(W, b) = \lambda \sum_{l=0}^2 \|W^{(l)}\|_2^2$$

where

$\lambda =$

$10^{-4}$ .

Figure

(??)

shows

the

pos-

te-

rior

dis-

tri-

bu-

tion

of

the

clas-

si-

fi-

ca-

tion

ac-

cu-

racy

in

the

test

dataset.

We

see

19720??

that

the

test

ac-

cu-

ra-

cies

are

cen-

tered

around

0.75

and

the

pos-

te-

rior

dis-

tri-

bu-

tion

is  
left-  
skewed.  
Fur-  
ther-  
more,  
the  
ac-  
cu-  
racy  
is  
higher  
than  
0.35  
in  
99%  
of  
the  
cases.  
The  
95%

19920??

in-  
ter-  
4.  
val-  
DIS-  
GUS-  
SION  
[0.407,  
Weighted  
Bayesian  
Boot-  
strap  
(WBB)  
pro-  
vides  
a  
com-  
pu-  
ta-  
tion-  
ally  
at-  
trac-  
tive  
so-

FIGURE 4:  
Posterior  
dis-  
tri-  
bu-  
tion  
of  
the  
clas-  
si-  
fi-  
ca-  
tion  
ac-  
cu-  
racy.  
 $n =$   
500,  $\lambda =$   
 $10^{-4}$ .

lu-  
tion  
to  
scal-  
able  
Bayesian  
in-  
fer-  
ence  
(??)  
whilst  
ac-  
count-  
ing  
for  
pa-  
ram-  
e-  
ter  
un-  
cer-



20120??

tainty

by

draw-

ing

sam-

ples

from

a

weighted

pos-

te-

rior

dis-

tri-

bu-

tion.

WBB

can

also

be

used

in

con-

junc-

tion

with

prox-

i-

mal

meth-

ods

(?,

?)

to

pro-

vide

spar-

sity

in

high

di-

men-

20320??

sional

sta-

tis-

tica

prob-

lems.

With

a

sim-

i-

lar

ease

of

com-

pu-

ta-

tion,

WBB

pro-

vides

an

al-

ter-

na-

tive

to

ABC

meth-

ods

(?)

and

Vari-

a-

tional

Bayes

(VB)

meth-

ods.

A

fruit-

ful

area

20520??

for

fu-

ture

re-

search

is

the

com-

par-

i-

son

of

ap-

prox-

i-

mate

Bayesian

com-

pu-

ta-

tion

with  
sim-  
u-  
lated  
Bayesian  
Boot-  
strap  
in-  
fer-  
ence.

## ACKNOWLEDGEMENTS

Place  
all  
ac-  
knowl-  
edge-  
ments  
here.  
In  
your  
ini-

20720??

tial

and

re-

vised

sub-

mis-

sion,

en-

sure

that

any

ac-

knowl-

edge-

ments

are

anony-

mous;

in-

clude

the

full  
ac-  
knowl-  
edge-  
ments  
only  
af-  
ter  
your  
pa-  
per  
has  
been  
ac-  
cepted.  
Grant-  
ing  
agen-  
cies  
should  
not



20920??

be

ab-

bre-

vi-

ated,

and

do

not

in-

clude

grant

num-

bers.

We

are

grate-

ful

for

your

as-

sis-

BLUNDEDA AND BLUNDEDA  
BLUNDEDA AND BLUNDEDA

tance

with

our

pub-

li-

ca-

tion

pro-

cess.

APPENDIX

AP-

PENDIXSTOCHASTIC

GRA-

DI-

ENT

DE-

SCENT

(SGD)

Stochastic

gra-

*The Canadian Journal of Statistics / La revue canadienne de statistique*DOI:

21120??

di-

ent

de-

scent

(SGD)

method

or

its

vari-

a-

tion

is

typ-

i-

cally

used

to

find

the

deep

learn-

ing

model

weights

by

min-

i-

miz-

ing

the

pe-

nal-

ized

loss

func-

tion,

$$\sum_{i=1}^n w_i l_i(y_i; \theta) +$$

$$\lambda w_p \phi(\theta).$$

The

method

min-

i-

21320??

mizes

the

func-

tion

by

tak-

ing

a

neg-

a-

tive

step

along

an

es-

ti-

mate

$g^k$

of

the

gra-

di-

ent

$$\nabla \left[ \sum_{i=1}^n w_i l_i(y_i; \theta^k) + \lambda w_p \phi(\theta^k) \right]$$

at

it-

er-

a-

tion

$k$ .

The

ap-

prox-

i-

mate

gra-

di-

ent

is

es-

ti-

mated

21520??

by

cal-

cu-

lat-

ing

$$g^k = \frac{n}{b_k} \sum_{i \in E_k} w_i \nabla l_i(y_i; \theta^k) + \lambda w_p \frac{n}{b_k} \nabla \phi(\theta^k)$$

Where

$$E_k \subset$$

$$\{1, \dots, n\}$$

and

$$b_k =$$

$$|E_k|$$

is

the

num-

ber

of

el-

e-

ments

in

$E_k$ .

When

$b_k >$

1

the

al-

go-

rithm

is

called

batch

SGD

and

sim-

ply

SGD

oth-

er-

wise.

A



21720??

usual

strat-

egy

to

choose

sub-

set

$E$

is

to

go

cycli-

cally

and

pick

con-

sec-

u-

tive

el-

e-

ments

of

$\{1, \dots, T\}$ ,

$E_{k+1} =$

$[E_k$

$\bmod n] +$

1.

The

ap-

prox-

i-

mated

di-

rec-

tion

$g^k$

is

cal-

cu-

lated

us-

21920??

ing

a

chain

rule

(aka

back-

propagation)

for

deep

learn-

ing.

It

is

an

un-

bi-

ased

es-

ti-

ma-

tor.

Thus,

at

each

it-

er-

a-

tion,

the

SGD

up-

dates

the

so-

lu-

tion

$$\theta^{k+1} = \theta^k - t_k g^k$$

For

deep

learn-

ing

22120??

ap-

pli-

ca-

tions

the

step

size

$t_k$

(a.k.a

learn-

ing

rate)

is

usu-

ally

kept

con-

stant

or

some

sim-

ple

step

size

re-

duc-

tion

strat-

egy

is

used,

$t_k =$

$a \exp(-kt)$ .

Ap-

pro-

pri-

ate

learn-

ing

rates

or

the

hy-

22320??

per-

pa-

ram-

e-

ters

of

re-

duc-

tion

sched-

ule

are

usu-

ally

found

em-

pir-

i-

cally

from

nu-

mer-

i-  
cal  
ex-  
per-  
i-  
ments  
and  
ob-  
ser-  
va-  
tions  
of  
the  
loss  
func-  
tion  
pro-  
gres-  
sion.



22520??

*Received 9*

*Au-*

*gust*

*2018*

*Accepted 8*

*Septem-*

*ber*

*2018*