Report on the revision of "Weighted Bayesian Bootstrap for Scalable Posterior Distributions"

I would like to thank the authors for responding to my comments and their work in revising the manuscript.

I think the ideas and contents presented in the manuscript deserve publication, though there is one major issue in the newly added Bayesian lasso example (see below). Also, I really think the paper needs to be more explicit about describing what the proposed method does and what it does not. The revision does not address my suspicion that WBB uncertainty estimate in general does *not* coincide with the posterior distribution. Additionally, no further information (such as an actual computing time) seems to have been added to quantify the claims of "fast and scalable." As a reader, I would like to obtain a reasonably accurate picture of what proposed method achieves from the title, abstract, and introduction without having to carefully go through the results.

The model selection discrepancy metric in the Bayesian lasso example makes no sense. An actual posterior distribution of Bayesian lasso places no positive mass at a coefficient being precisely at zero. (Incidentally, the fact that WBB samples can yield precise zeros with positive probability only illustrates that they do not coincide with samples from the posterior.)